

OSS Deployment Cost and Latency

Measured public deployments for Qwen/Qwen2.5-0.5B-Instruct. HF CPU is free but slow; Modal T4 gives the best interactive latency for the assignment demo.

OSS Deployment Average Latency



OSS Deployment Throughput



Deployment	Cost Model	Avg Latency	Median Latency	Avg Tokens/sec
Hugging Face CPU	\$0/month CPU Basic	33.81s	24.97s	0.79
Modal CPU	Usage-based CPU + memory	7.41s	5.87s	3.38
Modal T4 GPU	\$0.000164/sec GPU time + CPU/memory	1.61s	0.42s	30.32

Recommendation: keep Hugging Face Spaces as the zero-cost public fallback and use Modal T4 for demo-quality latency. For production, validate cost in the Modal dashboard because idle warm containers and CPU/memory charges also apply.