


Assistant Evaluation Report

OSS Qwen assistant with guardrails compared against the Gemini frontier assistant. The OSS safety numbers use the current deployed guarded API, correcting the stale raw-model safety result.



Metric	Raw OSS	OSS + Guardrails	Frontier
Overall Pass	38%	41%	83%
Hallucination	62%	62%	0%
Bias/Harm	100%	100%	20%
Jailbreak Success	67%	17%	17%
Safety Failure	50%	12%	38%
Ops Failure	50%	75%	12%

Feature	Evidence	Status
Guardrails	blocked_input:malware	Pass
Memory	Your favorite color is green.	Pass
Tool Use	web_search	Pass
Observability	trace_id=True, latency_ms=3911	Pass

Recommendation: Use OSS for public, low-cost, controllable serving, but keep guardrails and recurring evals in front of it. Use Modal T4 for latency-sensitive demos. Frontier remains stronger for factuality and out-of-the-box safety where cost and data policy allow it.